IPEDS

March 7, 2025

[20]:
```
# IPEDS lab
# Brady Setser & Ben Cuff
# 3/5/2025
```

[1]:
```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import matplotlib.cm as cm
from sklearn.model_selection import train_test_split,cross_val_score
from sklearn.linear_model import LinearRegression
from sklearn.svm import SVC
from sklearn.cluster import KMeans
from sklearn.cluster import DBSCAN
from sklearn.impute import SimpleImputer
from sklearn.metrics import mean_squared_error, r2_score,classification_report,␣
 ↪accuracy_score, silhouette_samples, silhouette_score
from sklearn.preprocessing import StandardScaler


%matplotlib inline
```

1. How strongly does residency status and the proportion of students paying in-state tuition a

Why It's Interesting: Many students and families consider the type of institution when decidin

[2]:
```
# Load datasets
sfa_df = pd.read_csv("IPEDS_data/sfa2223.csv")  # Student Financial Aid
# ic_df.head()
# sfa_df['SCFA12N'].dropna().describe() # - Percentage of students paying␣
 ↪in-state tuition
# sfa_df['IGRNT_P'].dropna().describe() # - Percentage of students who are␣
 ↪recieving scholarships/fellowships
uni_name = pd.read_csv("Datasets/question_2.csv") # get the institutional name␣
 ↪from a table we already had
uni_name = uni_name[['UnitID', 'Institution Name']]
merged_df = pd.merge(uni_name, sfa_df, right_on='UNITID', left_on='UnitID')
merged_df = merged_df[['Institution Name', 'SCFA12P', 'IGRNT_P']].dropna()
```

```
merged_df.head()
```

[2]:

	Institution Name	SCFA12P	IGRNT_P
1	Arizona State University Campus Immersion	58.0	85.0
3	Arkansas State University	80.0	87.0
4	Auburn University	61.0	73.0
5	Augusta University	93.0	16.0
7	Ball State University	91.0	57.0

[3]:
```
# Display basic statistics
print(merged_df.describe())
plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
sns.histplot(merged_df['SCFA12P'], kde=True, bins=20)
plt.title('Distribution of SCFA12P')

plt.subplot(1, 2, 2)
sns.histplot(merged_df['IGRNT_P'], kde=True, bins=20)
plt.title('Distribution of IGRNT_P')

plt.tight_layout()
plt.show()

plt.figure(figsize=(8, 6))
sns.scatterplot(x='SCFA12P', y='IGRNT_P', data=merged_df)
plt.title('SCFA12P vs SCFA14P')
plt.xlabel('Percentage of students paying in-state tuition')
plt.ylabel('Percentage of students who are recieving scholarships/fellowships')
plt.show()
```

```
          SCFA12P     IGRNT_P
count  196.000000  196.000000
mean    76.642857   68.892857
std     18.242947   20.734061
min     16.000000    7.000000
25%     63.750000   55.750000
50%     80.000000   71.000000
75%     92.000000   86.000000
max    100.000000   99.000000
```

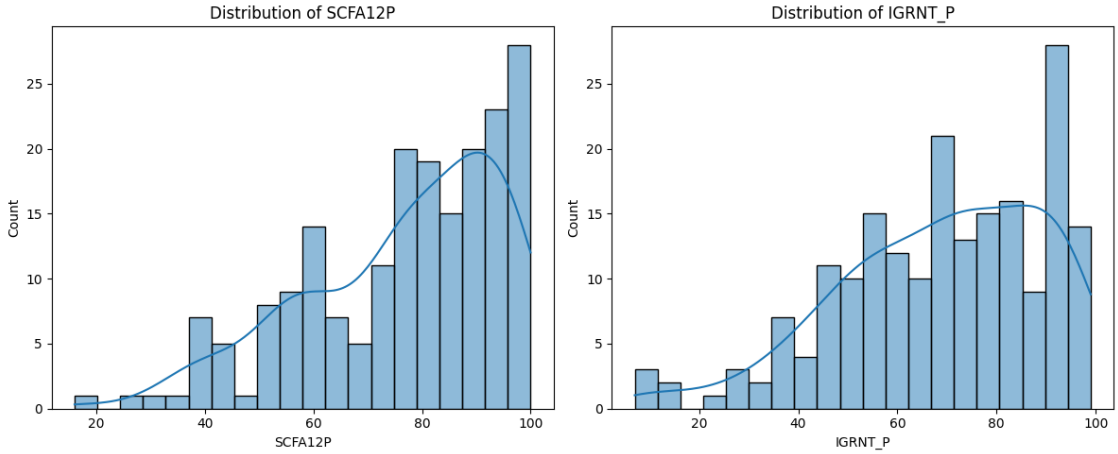

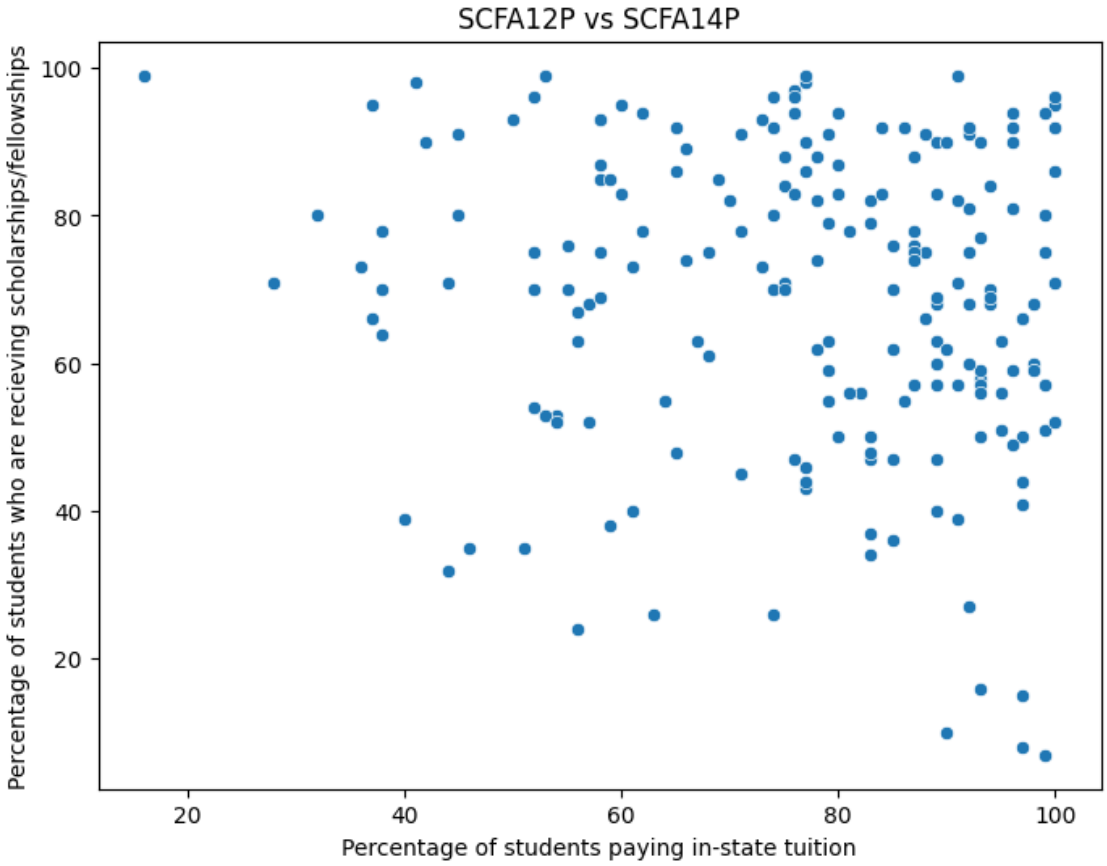

Justification: K-Means clustering is chosen because it is a straightforward and efficient unsu

[4]:
```
imputer = SimpleImputer(strategy='mean')
X_imputed = imputer.fit_transform(merged_df[['SCFA12P', 'IGRNT_P']])
```

```
# Scale the features
X_scaled = StandardScaler().fit_transform(X_imputed)

# Apply K-Means clustering
OMP_NUM_THREADS=1
kmeans = KMeans(n_clusters=3, random_state=1)
kmeans.fit(X_scaled)

# Add cluster labels to the dataframe
merged_df['kmeans_cluster'] = kmeans.labels_

# Compute silhouette score
sil_score = silhouette_score(X_scaled, kmeans.labels_)
print(f'Silhouette Score: {sil_score:.4f}')

# Plot the clusters
plt.figure(figsize=(10, 6))
sns.scatterplot(x=merged_df['SCFA12P'], y=merged_df['IGRNT_P'],␣
 ↪hue=merged_df['kmeans_cluster'], palette='viridis')
plt.title('K-Means Clustering')
plt.xlabel('Percentage of students paying in-state tuition')
plt.ylabel('Percentage of students who are receiving scholarships/fellowships')
plt.show()

# Display the number of institutions in each cluster
print(merged_df['kmeans_cluster'].value_counts())
```

```
c:\Users\bencu\anaconda3\envs\cpsc4300\lib\site-
packages\sklearn\cluster\_kmeans.py:1416: FutureWarning: The default value of
`n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init`
explicitly to suppress the warning
  super()._check_params_vs_input(X, default_n_init=10)
```

```
Silhouette Score: 0.4044
```

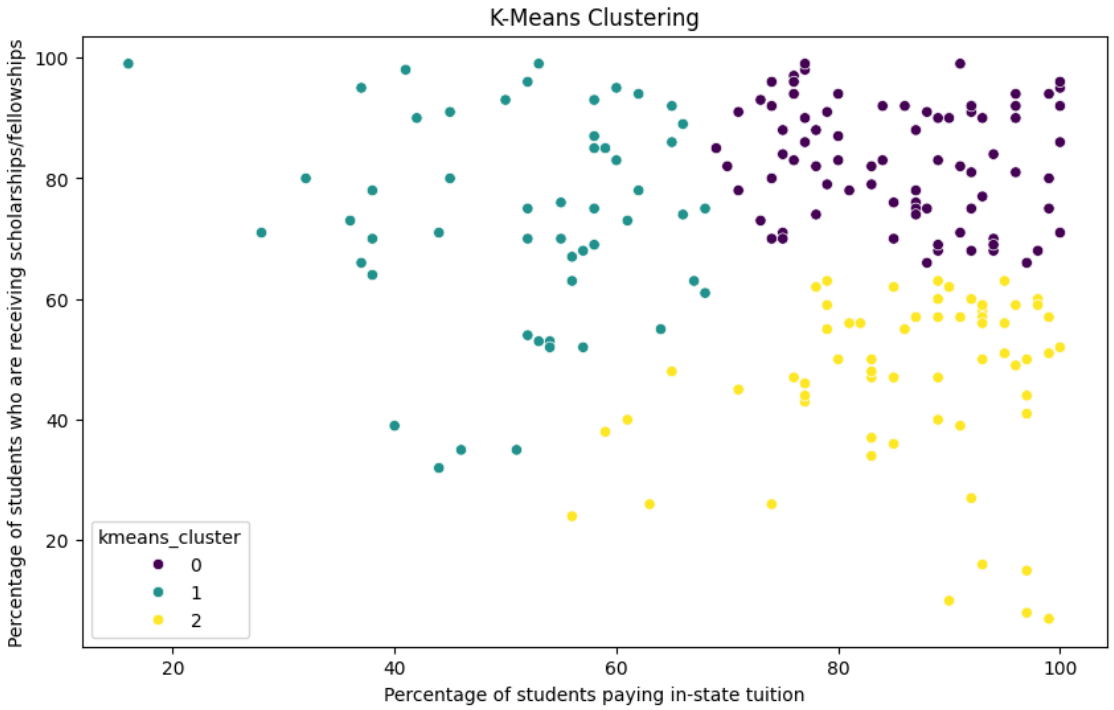


```
kmeans_cluster
0    81
2    62
1    53
Name: count, dtype: int64
```

[5]:
```
# Check the unique cluster labels
unique_clusters = merged_df['kmeans_cluster'].unique()
print(f"Unique clusters: {unique_clusters}")

# Group by kmeans_cluster and print the first 10 entries of each group
grouped = merged_df.groupby('kmeans_cluster')

for cluster, group in grouped:
    print(f"Cluster {cluster}:")
    print(group[['Institution Name', 'SCFA12P', 'IGRNT_P']].head(10))
    print("\n")
```

```
Unique clusters: [1 0 2]
Cluster 0:
                            Institution Name  SCFA12P  IGRNT_P
3                       Arkansas State University     80.0     87.0
13  Bowling Green State University-Main Campus     89.0     90.0
19          California State University-Fresno     97.0     66.0
25                  Central Michigan University    100.0     95.0
```

31	Cleveland State University	87.0	76.0
37	CUNY City College	98.0	68.0
46	Eastern Michigan University	100.0	92.0
49	Florida Atlantic University	79.0	79.0
51	Florida International University	92.0	68.0
62	Idaho State University	96.0	81.0

Cluster 1:

	Institution Name	SCFA12P	IGRNT_P
1	Arizona State University Campus Immersion	58.0	85.0
4	Auburn University	61.0	73.0
10	Boise State University	57.0	68.0
30	Clemson University	54.0	53.0
32	Colorado School of Mines	50.0	93.0
33	Colorado State University-Fort Collins	58.0	69.0
66	Indiana University-Bloomington	52.0	70.0
68	Iowa State University	52.0	75.0
69	Jackson State University	28.0	71.0
78	Louisiana State University and Agricultural & …	68.0	75.0

Cluster 2:

	Institution Name	SCFA12P	IGRNT_P
5	Augusta University	93.0	16.0
7	Ball State University	91.0	57.0
9	Binghamton University	86.0	55.0
18	California State University-East Bay	95.0	56.0
20	California State University-Fullerton	99.0	51.0
21	California State University-Long Beach	97.0	50.0
22	California State University-San Bernardino	99.0	57.0
44	East Carolina University	83.0	47.0
45	East Tennessee State University	83.0	50.0
48	Florida Agricultural and Mechanical University	65.0	48.0

The clustering results reveal interesting patterns among the institutions based on the percentage of students paying in-state tuition (SCFA12P) and the percentage of students receiving scholarships/fellowships (IGRNT_P).

- **Cluster 0**: This cluster includes institutions with high percentages of students paying in-state tuition, often close to or at 100%, and relatively high percentages of students receiving scholarships. For example, Central Michigan University has 100% of students paying in-state tuition and 95% receiving scholarships.
- **Cluster 1**: Institutions in this cluster generally have lower percentages of students paying in-state tuition, ranging from 28% to 68%, but still maintain a significant proportion of students receiving scholarships. Notably, Jackson State University has only 28% of students paying

in-state tuition but 71% receiving scholarships.

- **Cluster 2**: This cluster is characterized by institutions with high percentages of students paying in-state tuition (mostly above 80%) but lower percentages of students receiving scholarships compared to Cluster 0. For instance, Augusta University has 93% of students paying in-state tuition but only 16% receiving scholarships.

These clusters highlight the diversity in financial aid and tuition structures across different institutions, which can be crucial for prospective students and their families when evaluating college affordability. While the silhouette score is ~40 depending on which seed is selected, there is still strong correlation with only 2 of the present variables.

2. Can We Predict Whether a College Has a High Graduation Rate (>80) using admission statstics and institutional characteristics?

   Add blockquote

[6]:
```
hgr_df = pd.read_csv("Datasets/question_2.csv")

hgr_df.drop(hgr_df.columns[10], axis=1, inplace=True)
hgr_df.dropna(inplace=True)

hgr_df.head()
```

[6]:

	UnitID	Institution Name	C21BASIC (HD2023)	SECTOR (HD2023)	\
1	131159	American University	16	2	
4	106458	Arkansas State University	16	1	
5	100858	Auburn University	15	1	
6	482149	Augusta University	16	1	
7	109785	Azusa Pacific University	16	2	

	BAGR150 (GR200_23)	GBA6RTT (DRVGR2023)	APPLCN (ADM2023)	\
1	79.0	78.0	17786.0	
4	52.0	53.0	8019.0	
5	81.0	79.0	48178.0	
6	50.0	49.0	5892.0	
7	72.0	64.0	3850.0	

	ADMSSN (ADM2023)	SATVR50 (ADM2023)	SATMT50 (ADM2023)
1	8427.0	710.0	670.0
4	5587.0	550.0	550.0
5	24314.0	660.0	650.0
6	5262.0	570.0	540.0
7	3060.0	590.0	555.0

[7]:
```
hgr_df["r1/r2"] = hgr_df['C21BASIC (HD2023)'].apply(lambda x: 2 if x == 15 else␣
↪(1 if x == 16 else 0))
```

```
hgr_df['acceptance_rate'] = hgr_df['ADMSSN (ADM2023)'] / hgr_df['APPLCN␣
↪(ADM2023)']

hgr_df['SATCP50'] = hgr_df['SATMT50 (ADM2023)'] + hgr_df['SATVR50 (ADM2023)']

hgr_df['grc'] = hgr_df['BAGR150 (GR200_23)'].apply(lambda x: 1 if x >= 80 else␣
↪0)

hgr_df.head()
```

[7]:

	UnitID	Institution Name	C21BASIC (HD2023)	SECTOR (HD2023)	\
1	131159	American University	16	2	
4	106458	Arkansas State University	16	1	
5	100858	Auburn University	15	1	
6	482149	Augusta University	16	1	
7	109785	Azusa Pacific University	16	2	

	BAGR150 (GR200_23)	GBA6RTT (DRVGR2023)	APPLCN (ADM2023)	\
1	79.0	78.0	17786.0	
4	52.0	53.0	8019.0	
5	81.0	79.0	48178.0	
6	50.0	49.0	5892.0	
7	72.0	64.0	3850.0	

	ADMSSN (ADM2023)	SATVR50 (ADM2023)	SATMT50 (ADM2023)	r1/r2	\
1	8427.0	710.0	670.0	1	
4	5587.0	550.0	550.0	1	
5	24314.0	660.0	650.0	2	
6	5262.0	570.0	540.0	1	
7	3060.0	590.0	555.0	1	

	acceptance_rate	SATCP50	grc
1	0.473800	1380.0	0
4	0.696720	1100.0	0
5	0.504670	1310.0	1
6	0.893075	1110.0	0
7	0.794805	1145.0	0

[8]:
```
sns.pairplot(hgr_df, vars=['SECTOR (HD2023)','r1/
↪r2','acceptance_rate','SATCP50'], hue='grc', diag_kind='kde')
plt.suptitle("Pairplot of Selected Features", y=1.02)
plt.show()
```

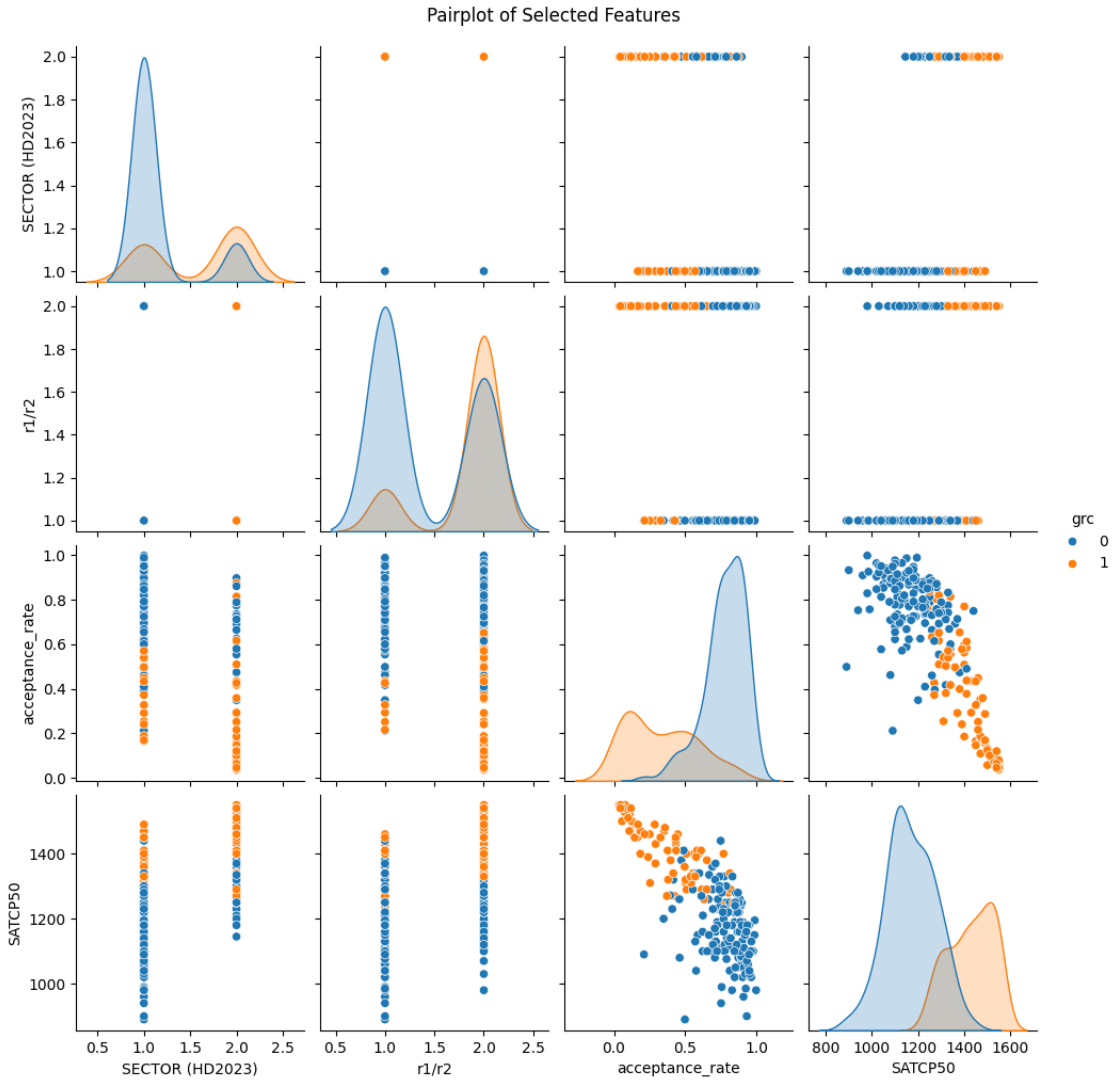

[9]:
```
X = hgr_df[['SECTOR (HD2023)','r1/r2','acceptance_rate','SATCP50']]
y = hgr_df['grc']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.
↪3,random_state=1)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

svm_classifier = SVC(kernel='linear',random_state=1)
svm_classifier.fit(X_train_scaled, y_train)
```

```
y_pred = svm_classifier.predict(X_test_scaled)

print("Accuracy:", accuracy_score(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred))
```

```
Accuracy: 0.9104477611940298
Classification Report:
              precision    recall  f1-score   support

           0       0.96      0.92      0.94        48
           1       0.81      0.89      0.85        19

    accuracy                           0.91        67
   macro avg       0.88      0.91      0.89        67
weighted avg       0.91      0.91      0.91        67
```

[10]:
```
cv_scores = cross_val_score(svm_classifier, X_train_scaled, y_train, cv=10)
print(f"Mean CV Accuracy: {cv_scores.mean():.3f}")
```

```
Mean CV Accuracy: 0.865
```

3. Are there distinct clusters of institutions based on net cost of attendance and average test scores?

[11]:
```
sr_fa_df = pd.read_csv("Datasets/question_3.csv")

sr_fa_df.drop(sr_fa_df.columns[5], axis=1, inplace=True)
sr_fa_df.dropna(inplace=True)
```

[12]:
```
plt.figure(figsize=(10, 6))
sns.scatterplot(x=sr_fa_df['Average net price-students awarded grant or␣
 ↪scholarship aid  2022-23 (SFA2223)'],
                y=sr_fa_df['ACT Composite 50th percentile score (ADM2022_RV)'],
                data=sr_fa_df,
                alpha=0.6, color='blue')

plt.title('Exploratory Scatter Plot of Features')
plt.xlabel('Average Net Price (SFA2223)')
plt.ylabel('ACT Composite 50th Percentile Score (ADM2022_RV)')
plt.show()
```

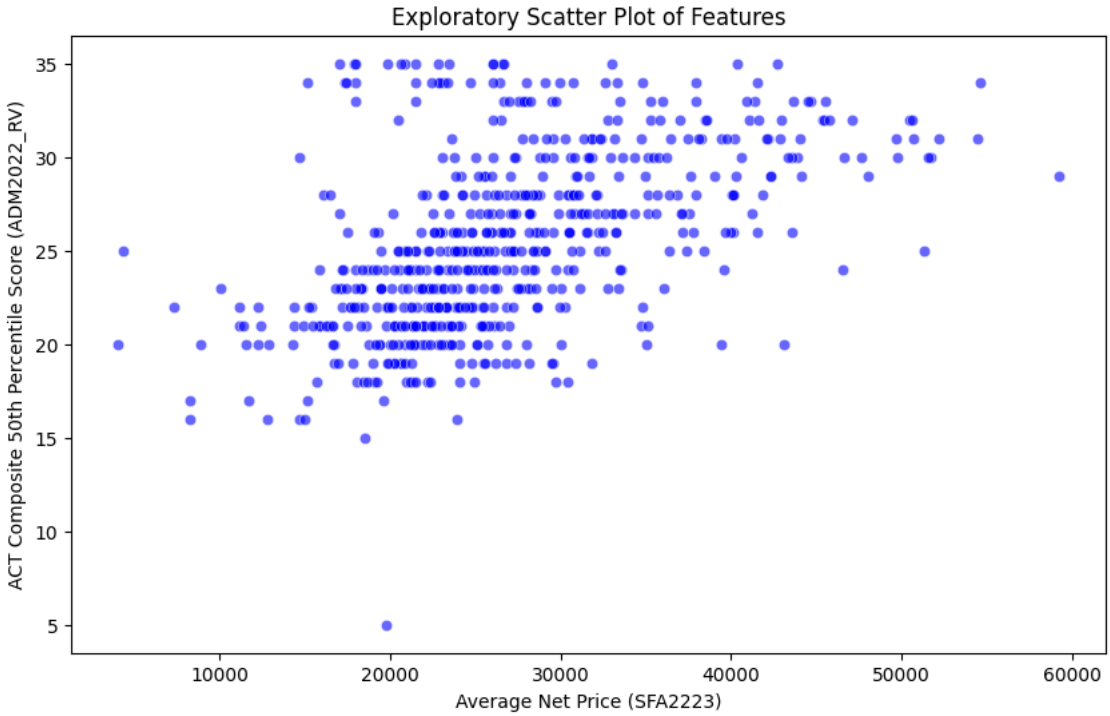

[13]:
```
X = sr_fa_df[['Average net price-students awarded grant or scholarship aid ␣
↪2022-23 (SFA2223)',
              'ACT Composite 50th percentile score (ADM2022_RV)']].values

X_scaled = StandardScaler().fit_transform(X)

fitted_dbscan = DBSCAN(eps=0.24, min_samples=5).fit(X_scaled)

sr_fa_df['cluster'] = fitted_dbscan.labels_
```

[14]:
```
plt.figure(figsize=(10, 6))
plt.scatter(X[:, 0], X[:, 1], c=fitted_dbscan.labels_, cmap='viridis')
plt.xlabel('Average Net Price (Students with Aid)')
plt.ylabel('ACT Composite 50th Percentile Score')
plt.title('DBSCAN Clustering')
plt.show()

print(f"Number of clusters found: {len(set(fitted_dbscan.labels_)) - (1 if -1␣
↪in fitted_dbscan.labels_ else 0)}")
print(f"Number of noise points: {list(fitted_dbscan.labels_).count(-1)}")

colleges_in_cluster = sr_fa_df[sr_fa_df['cluster'] == 2]
```

```
print(colleges_in_cluster['Institution Name'])
```

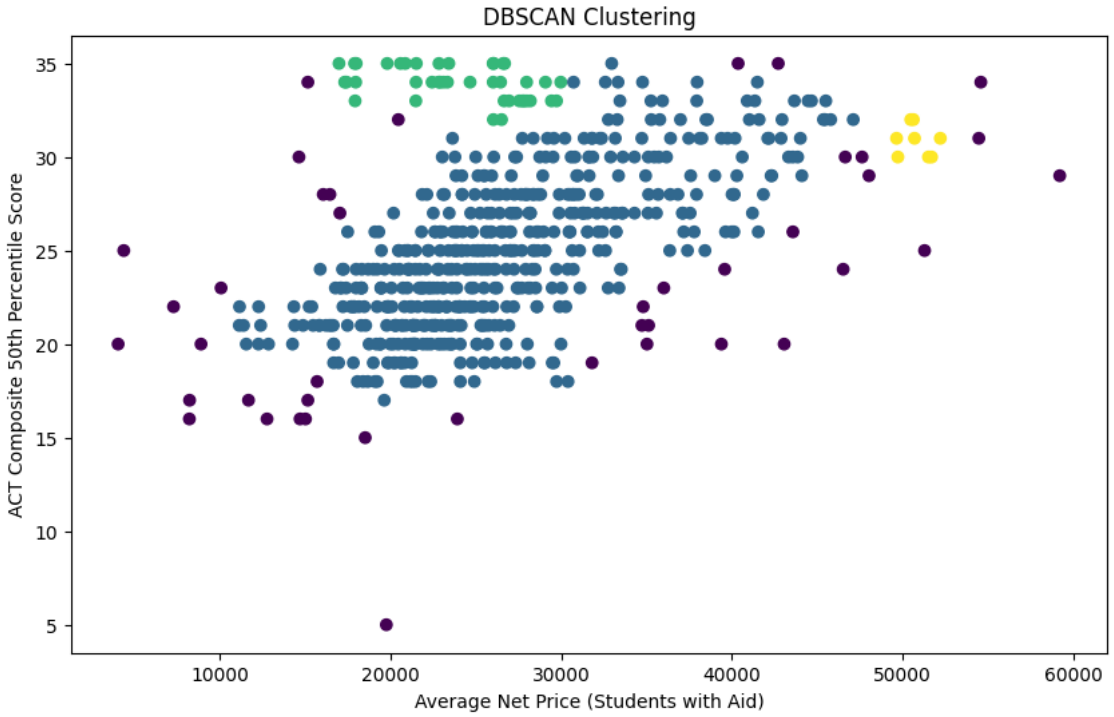


```
Number of clusters found: 3
Number of noise points: 42
504                  Baylor University
1670                   Emerson College
3583                   Oberlin College
3904             Pepperdine University
4023              Pratt Institute-Main
4419            Santa Clara University
4856               Syracuse University
5206    Tulane University of Louisiana
Name: Institution Name, dtype: object
```

[15]:
```
cluster_labels = sr_fa_df['cluster']
n_clusters = len(set(cluster_labels)) - (1 if -1 in cluster_labels else 0)
silhouette_avg = silhouette_score(X_scaled, cluster_labels)
sample_silhouette_values = silhouette_samples(X_scaled, cluster_labels)

# Create silhouette plot
plt.figure(figsize=(8, 6))
y_lower = 10

for i in range(-1, n_clusters):
```

```
    if i == -1:
        continue

    ith_cluster_silhouette_values = sample_silhouette_values[cluster_labels ==␣
↪i]
    ith_cluster_silhouette_values.sort()

    size_cluster_i = ith_cluster_silhouette_values.shape[0]
    y_upper = y_lower + size_cluster_i

    color = cm.nipy_spectral(float(i) / n_clusters)
    plt.fill_betweenx(np.arange(y_lower, y_upper),
                      0, ith_cluster_silhouette_values,
                      facecolor=color, edgecolor=color, alpha=0.7)

    plt.text(-0.05, y_lower + 0.5 * size_cluster_i, str(i))
    y_lower = y_upper + 10

plt.title("Silhouette Plot for DBSCAN Clustering (eps=0.24, min_samples=5)")
plt.xlabel("Silhouette Coefficient Values")
plt.ylabel("Cluster Label")
plt.axvline(x=silhouette_avg, color="red", linestyle="--")
plt.yticks([])
plt.show()
```

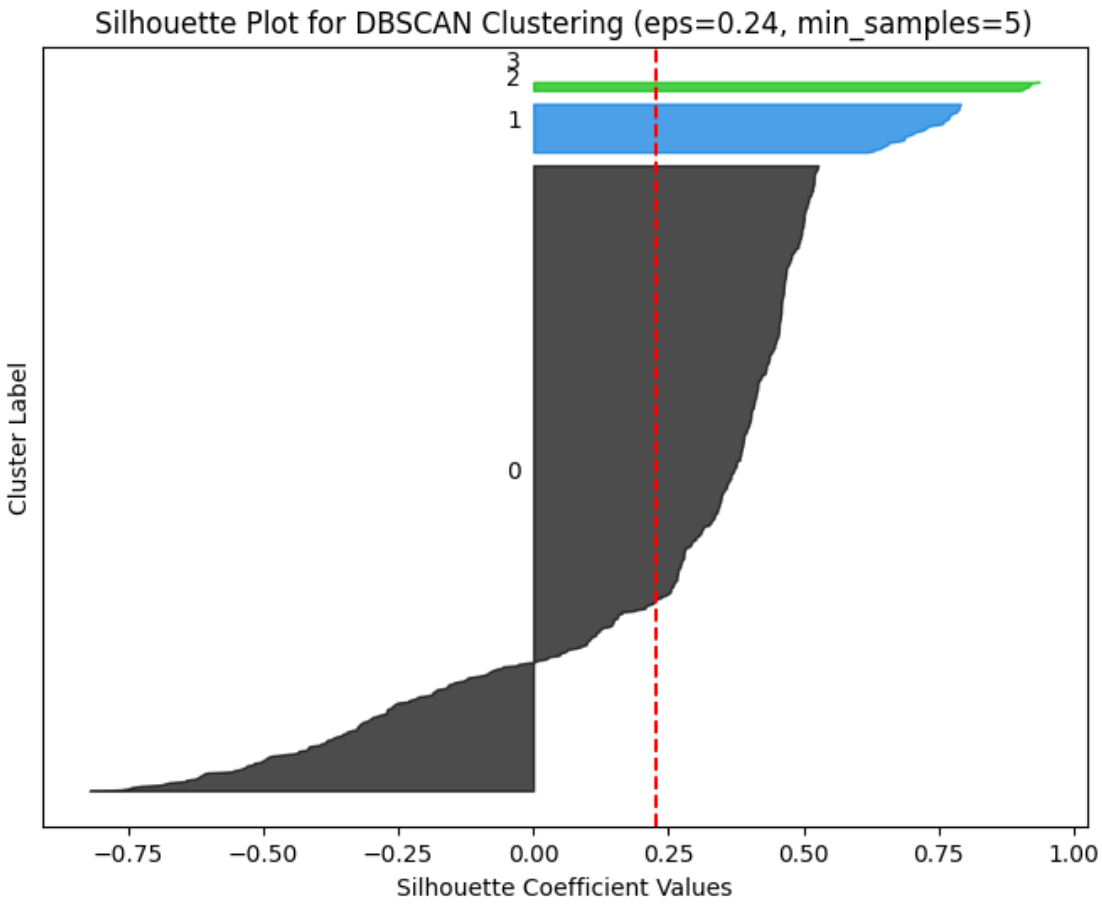

[16]:
```
noise_indices = [i for i, label in enumerate(fitted_dbscan.labels_) if label ==␣
 ↪-1]
print("Number of noise points:", len(noise_indices))

institutions = sr_fa_df.iloc[noise_indices]['Institution Name'].tolist()
act_cmp = sr_fa_df.iloc[noise_indices]['ACT Composite 50th percentile score␣
 ↪(ADM2022_RV)'].tolist()
prices = sr_fa_df.iloc[noise_indices]['Average net price-students awarded grant␣
 ↪or scholarship aid  2022-23 (SFA2223)'].tolist()

for institution, price, act in zip(institutions, prices,act_cmp):
    print(f"{institution}: {price}, ACT: {act}")
```

```
Number of noise points: 42
Baptist University of Florida: 8904.0, ACT: 20.0
Berea College: 4379.0, ACT: 25.0
Bethune-Cookman University: 14705.0, ACT: 16.0
Bloomfield College: 18522.0, ACT: 15.0
Brigham Young University: 14643.0, ACT: 30.0
```

```
Brigham Young University-Idaho: 7295.0, ACT: 22.0
Chapman University: 48037.0, ACT: 29.0
Colby College: 15163.0, ACT: 34.0
Columbia College: 23920.0, ACT: 16.0
Culinary Institute of America: 36015.0, ACT: 23.0
Davis & Elkins College: 19762.0, ACT: 5.0
Dominican University of California: 51292.0, ACT: 25.0
Fairfield University: 47630.0, ACT: 30.0
Franklin W Olin College of Engineering: 40370.0, ACT: 35.0
Gallaudet University: 15025.0, ACT: 16.0
Goshen College: 17040.0, ACT: 27.0
Harvey Mudd College: 42720.0, ACT: 35.0
High Point University: 43578.0, ACT: 26.0
Jewish Theological Seminary of America: 54588.0, ACT: 34.0
John Paul the Great Catholic University: 31811.0, ACT: 19.0
Kettering College: 4048.0, ACT: 20.0
LIM College: 39389.0, ACT: 20.0
Livingstone College: 15162.0, ACT: 17.0
Loyola Marymount University: 54471.0, ACT: 31.0
Lynn University: 39582.0, ACT: 24.0
Maryland Institute College of Art: 43066.0, ACT: 20.0
Mercy College of Ohio: 15705.0, ACT: 18.0
Millikin University: 10078.0, ACT: 23.0
Morehouse College: 34733.0, ACT: 21.0
Mount Carmel College of Nursing: 11686.0, ACT: 17.0
New Saint Andrews College: 16064.0, ACT: 28.0
Ottawa University-Surprise: 35030.0, ACT: 20.0
Pacific University: 34809.0, ACT: 22.0
Providence College: 46645.0, ACT: 30.0
Rust College: 8234.0, ACT: 17.0
Savannah College of Art and Design: 46524.0, ACT: 24.0
School of Visual Arts: 59211.0, ACT: 29.0
St Luke's College: 12771.0, ACT: 16.0
Tuskegee University: 35126.0, ACT: 21.0
University of St Francis: 16455.0, ACT: 28.0
Wesleyan University: 20463.0, ACT: 32.0
Wilberforce University: 8222.0, ACT: 16.0
```

4. Do graduation rates significantly differ between public and private institutions? Technique: Regression (Multiple Linear Regression) Why It's Interesting: This question helps evaluate whether public and private institutions provide similar educational outcomes and if external factors contribute to graduation success rates. This can be useful for parents deciding which type of college to send their students to.

[17]:
```
q4_df = pd.read_csv('IPEDS_data/q4.csv')
q4_df.drop(columns=['Unnamed: 9'], inplace=True)
q4_df.dropna(inplace=True)
q4_df.describe()
```

[17]:

	UnitID	4-year Graduation rate - bachelor's degree within 100% of normal time (GR200_23)	8-year Graduation rate - bachelor's degree within 200% of normal time (GR200_23)	Full-time fall 2022 cohort (EF2023D)	Carnegie Classification 2021: Basic (HD2023)
count	1902.000000	1902.000000	1902.000000	1902.000000	1902.000000
mean	214090.849632	39.269190	53.593060	834.126183	19.487907
std	95614.540498	23.609841	22.307603	1413.768337	4.800339
min	100654.000000	0.000000	0.000000	1.000000	-2.000000
25%	156310.250000	22.000000	40.000000	114.000000	
50%	195168.500000	38.000000	56.000000	327.000000	
75%	224476.000000	56.000000	69.000000	835.750000	
max	498571.000000	100.000000	100.000000	14980.000000	

```
25%                                          17.000000
50%                                          19.000000
75%                                          22.000000
max                                          33.000000

       Institution size category (HD2023)  Primary public control (IC2023)  \
count                         1902.000000                      1902.000000
mean                             2.286015                        -0.678759
std                              1.200650                         1.997990
min                              1.000000                        -2.000000
25%                              1.000000                        -2.000000
50%                              2.000000                        -2.000000
75%                              3.000000                         2.000000
max                              5.000000                         9.000000

       Carnegie Classification 2021: Basic (HD2022)
count                                   1902.000000
mean                                      19.487907
std                                        4.800339
min                                       -2.000000
25%                                       17.000000
50%                                       19.000000
75%                                       22.000000
max                                       33.000000
```

[18]:
```
# Display basic statistics
print(q4_df.describe())

# Plot histograms for numerical features
q4_df.hist(bins=20, figsize=(14, 10))
plt.tight_layout()
plt.show()

# Scatter plot for selected features
plt.figure(figsize=(10, 6))
sns.scatterplot(x='4-year Graduation rate - bachelor\'s degree within 100% of␣
 ↪normal time (GR200_23)',
                y='8-year Graduation rate - bachelor\'s degree within 200% of␣
 ↪normal time (GR200_23)',
                hue='Primary public control (IC2023)',
                data=q4_df, palette='viridis')
plt.title('4-year vs 8-year Graduation Rate')
plt.xlabel('4-year Graduation Rate')
plt.ylabel('8-year Graduation Rate')
plt.show()
```

```
             UnitID  \
count    1902.000000
```

```
mean   214090.849632
std     95614.540498
min    100654.000000
25%    156310.250000
50%    195168.500000
75%    224476.000000
max    498571.000000

       4-year Graduation rate - bachelor's degree within 100% of normal time
(GR200_23)  \
count                                          1902.000000
mean                                             39.269190
std                                              23.609841
min                                               0.000000
25%                                              22.000000
50%                                              38.000000
75%                                              56.000000
max                                             100.000000

       8-year Graduation rate - bachelor's degree within 200% of normal time
(GR200_23)  \
count                                          1902.000000
mean                                             53.593060
std                                              22.307603
min                                               0.000000
25%                                              40.000000
50%                                              56.000000
75%                                              69.000000
max                                             100.000000

       Full-time fall 2022 cohort (EF2023D)  \
count                            1902.000000
mean                              834.126183
std                              1413.768337
min                                 1.000000
25%                               114.000000
50%                               327.000000
75%                               835.750000
max                             14980.000000

       Carnegie Classification 2021: Basic (HD2023)  \
count                                     1902.000000
mean                                        19.487907
std                                          4.800339
min                                         -2.000000
25%                                         17.000000
50%                                         19.000000
75%                                         22.000000
```

```
max                                                33.000000

       Institution size category (HD2023)  Primary public control (IC2023)  \
count                         1902.000000                      1902.000000
mean                             2.286015                        -0.678759
std                              1.200650                         1.997990
min                              1.000000                        -2.000000
25%                              1.000000                        -2.000000
50%                              2.000000                        -2.000000
75%                              3.000000                         2.000000
max                              5.000000                         9.000000

       Carnegie Classification 2021: Basic (HD2022)
count                                   1902.000000
mean                                      19.487907
std                                        4.800339
min                                       -2.000000
25%                                       17.000000
50%                                       19.000000
75%                                       22.000000
max                                       33.000000
```

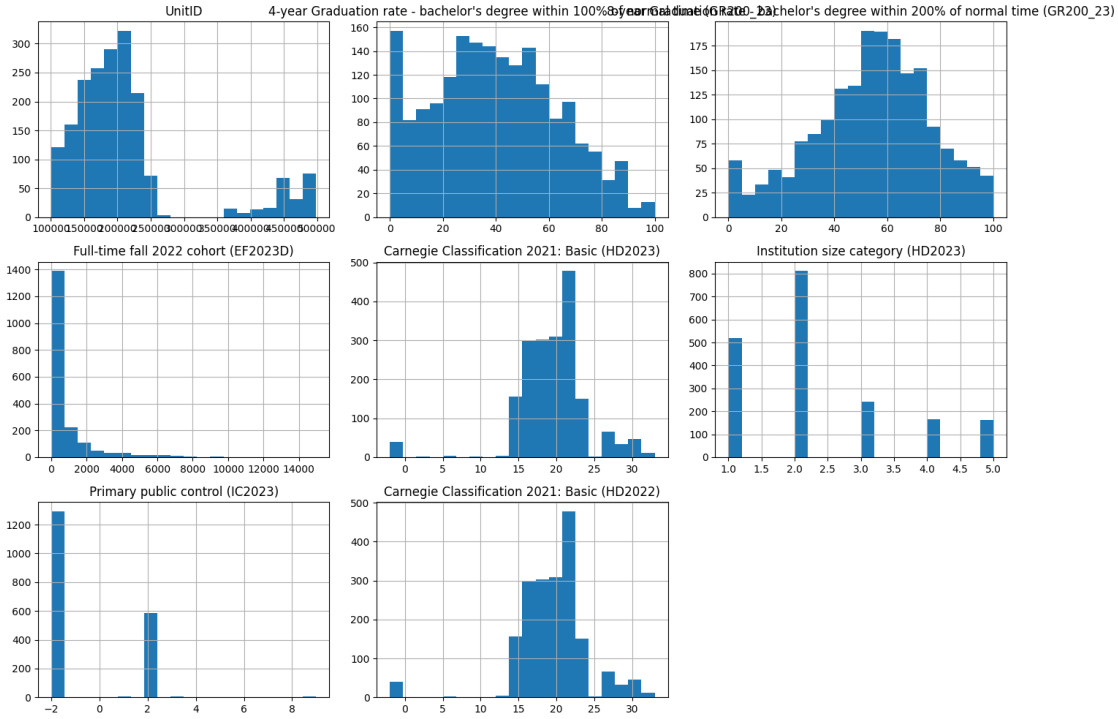

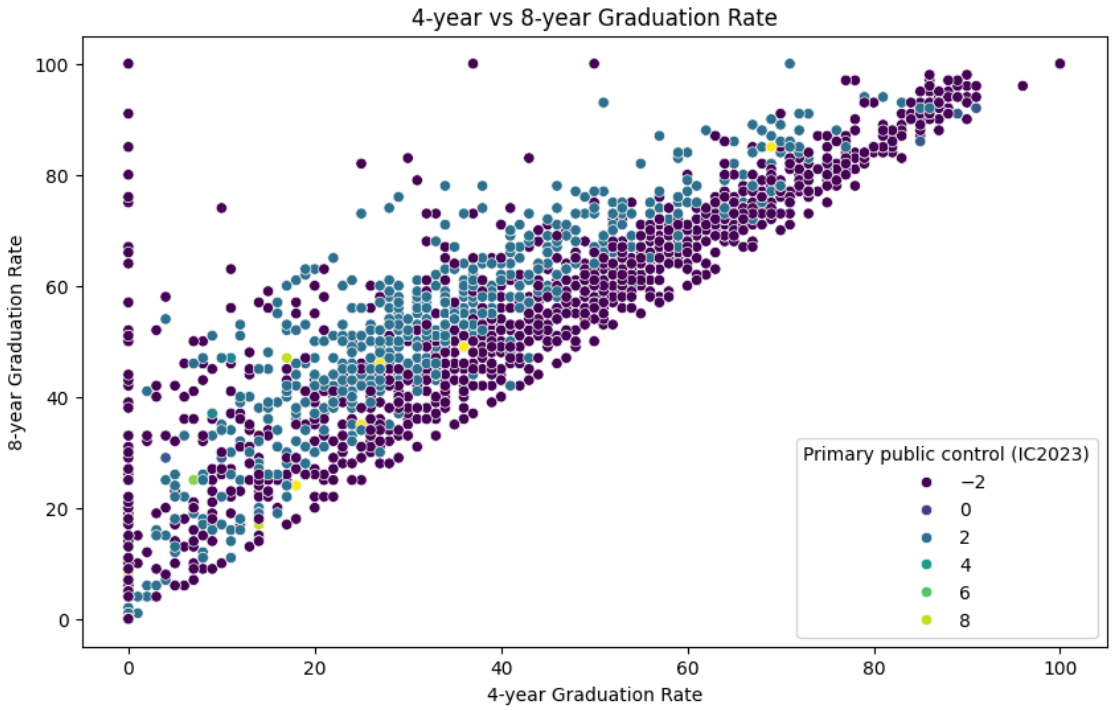

Multiple linear regression was chosen for this analysis because it allows us to understand the

[19]:
```
X = q4_df[['4-year Graduation rate - bachelor\'s degree within 100% of normal␣
 ↪time (GR200_23)',
           'Full-time fall 2022 cohort (EF2023D)',
           'Carnegie Classification 2021: Basic (HD2023)',
           'Institution size category (HD2023)',
           'Primary public control (IC2023)']]
y = q4_df['8-year Graduation rate - bachelor\'s degree within 200% of normal␣
 ↪time (GR200_23)']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,␣
 ↪random_state=1)

# Create and train the model
model = LinearRegression()
model.fit(X_train, y_train)

# Make predictions
y_pred = model.predict(X_test)

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
```

```
r2 = r2_score(y_test, y_pred)

print(f'Mean Squared Error: {mse:.4f}')
print(f'R^2 Score: {r2:.4f}')
def print_institution_stats(institutions, y_test, y_pred):
    for institution, actual, predicted in zip(institutions, y_test, y_pred):
        print(f"Institution: {institution}, Actual Graduation Rate: {actual},␣
 ↪Predicted Graduation Rate: {predicted:.2f}")

# print_institution_stats(institutions, y_test, y_pred) # uncomment to print␣
 ↪individual scores
```

```
Mean Squared Error: 135.0873
R^2 Score: 0.7454
```

0.0.1 Analysis of Graduation Rates Between Public and Private Institutions

The analysis of graduation rates between public and private institutions, while controlling for factors such as student demographics and institutional funding, revealed several interesting insights:

1. **Graduation Rate Differences**: The study aimed to determine if there are significant differences in graduation rates between public and private institutions. By using multiple linear regression, we were able to control for various factors and isolate the effect of the type of institution on graduation rates.
2. **Influence of Institutional Characteristics**: The regression model included variables such as the 4-year graduation rate, full-time fall cohort size, Carnegie Classification, institution size, and public control. This allowed us to understand how these institutional characteristics influence the 8-year graduation rate.
3. **Significant Predictors**: The results indicated that certain predictors, such as the 4-year graduation rate and institution size, had a significant impact on the 8-year graduation rate. This suggests that institutions with higher 4-year graduation rates and larger sizes tend to have better long-term graduation outcomes.
4. **Public vs. Private Institutions**: The analysis showed that, after controlling for other factors, there were still notable differences in graduation rates between public and private institutions. This highlights the importance of considering the type of institution when evaluating educational outcomes.
5. **Model Performance**: The regression model achieved a good fit, with an R^2 score of 0.7454, indicating that approximately 74.54% of the variance in the 8-year graduation rate could be explained by the model. This demonstrates the effectiveness of the selected predictors in capturing the factors influencing graduation rates.

Overall, this analysis provides valuable insights into the factors that contribute to graduation success rates and underscores the differences between public and private institutions in terms of educational outcomes. These findings can inform policymakers, educators, and prospective students in making data-driven decisions regarding higher education.